

# Training and Predicting with Linear Regression Models in scikit-learn

In this lecture, Professor Ashish explains how to train and make predictions using linear regression models in scikit-learn, covering both the normal equation method and iterative optimization using stochastic gradient descent (SGD).

## Training Linear Regression Models

The lecture details two approaches for training linear regression models:

1. **Normal Equation:** This method uses the `LinearRegression` class from `sklearn.linear_model`. The training process involves instantiating the `LinearRegression` object and then calling the `fit` method with the training feature matrix (`X_train`) and label vector (`y_train`).
2. **Iterative Optimization (SGD):** This approach employs the `SGDRegressor` class, suitable for larger datasets. The `fit` method is used similarly, but `SGDRegressor` offers more control over the optimization process through various parameters.

## SGDRegressor Parameters and Optimization

The lecture emphasizes the importance of several `SGDRegressor` parameters:

- **learning\_rate:** Controls the step size during optimization. Inverse scaling is often used, calculated as  $\eta_t = \frac{\eta_0}{t^{\text{power}_t}}$ , where  $\eta_0$  is the initial learning rate,  $t$  is the current iteration, and `power_t` is a hyperparameter.
- **average:** Enables averaging of the weights after a specified number of samples, improving performance with many features and a high  $\eta_0$ . Setting `average` to an integer ( $n$ ) starts averaging after seeing ( $n$ ) samples.
- **warm\_start = True:** This allows initializing SGD with the weight vector from a previous run, useful for monitoring loss iteration by iteration. Setting `max_iter = 1` and iteratively calling `fit` with `warm_start = True` enables this.

**Tip:** Always set a `random_state` in the `SGDRegressor` constructor for reproducible results.

## Making Predictions

Prediction on new data involves two steps:

1. Arrange the data into a feature matrix (shape (`#samples`, `#features`)) or a sparse matrix.
2. Call the `predict` method on the trained linear regression object (either `LinearRegression` or `SGDRegressor`) with the feature matrix as an argument. The method returns the predicted labels.

## Feature Scaling and DummyRegressor

The lecture also touches upon feature scaling, highlighting that SGD is sensitive to it. It recommends using `StandardScaler` for preprocessing. Finally, it briefly introduces `DummyRegressor`, a baseline model that predicts using strategies like mean, median, quantile, or a constant value.

Prof. Ashish concludes by summarizing the key steps in training and using linear regression models in scikit-learn, emphasizing the flexibility and control offered by `SGDRegressor` for large datasets and the importance of proper parameter tuning and feature scaling.