

Mock Test Report

Student: Amit Kulkarni

Date: April 01, 2025

15
out of 50

Suggestions to Improve

The student demonstrates some understanding of reinforcement learning and data preprocessing but needs significant improvement in several key areas.

Topic: Reinforcement Learning (Reward Processing) - Review the concept of discounted cumulative rewards and how they are calculated using discount factors. Practice implementing the `discount_rewards` function with various reward sequences and discount factors. Pay close attention to the iterative nature of the calculation, working backwards from the final reward.

Topic: Data Preprocessing for Machine Learning - Research and compare different techniques for handling categorical features with high cardinality (many unique values). Explore methods like one-hot encoding, target encoding, binary encoding, and embedding techniques. Consider the trade-offs of each method in terms of dimensionality and computational cost.

Topic: Handling Imbalanced Datasets - Study various techniques for addressing class imbalance in classification problems. Explore resampling methods (oversampling, undersampling, SMOTE), cost-sensitive learning, and ensemble methods. Compare their effectiveness and understand when each technique is most appropriate.

Topic: TensorFlow/Keras for Deep Learning - Review the basics of TensorFlow/Keras, focusing on building and using sequential models. Practice making predictions using a trained model and understand how to handle input shapes correctly when using `model(obs[np.newaxis])`. Understand the output shape of a Keras model's prediction.

Practice coding the solutions to the problems you missed. Focus on understanding the underlying logic and debugging your code.

Work through additional exercises and examples to reinforce your understanding of the concepts.

Consult relevant documentation and tutorials for clarification on specific functions and libraries.

Correct Answers

Question

Correct Answer

What will be the output of the following code? ```python
import numpy as np
rewards = [10, 0, -50]
discount_factor = 0.8
discounted_rewards = np.array(rewards)
for step in range(len(rewards) - 2, -1, -1):
 discounted_rewards[step] += discounted_rewards[step + 1] * discount_factor
print(discounted_rewards) ```

[10, -22, -50]

Which of the following best describes the purpose of the `discount_rewards` function (assuming `rewards` is a list of rewards and `discount_factor` is a float)? ```python
def discount_rewards(rewards, discount_factor):
 discounted = np.array(rewards)
 for step in range(len(rewards) - 2, -1, -1):
 discounted[step] += discounted[step + 1] * discount_factor
 return discounted ```

It discounts future rewards based on the discount factor.

The following code snippet is intended to normalize rewards. What is the flaw in this implementation? ```python
import numpy as np
def normalize_rewards(rewards):
 return (rewards - np.mean(rewards)) / np.std(rewards) ```

It doesn't handle the case where the standard deviation is zero.

Given the following code, what will `all_rewards` contain after execution? ```python
import numpy as np
def play_one_step(env, obs):
 #Simplified for brevity # ... some code ...
 return [10, 20, 30]
Example rewards
rewards_list = []
for i in range(2):
 rewards = play_one_step(None, None)
 #Simulating env
 rewards_list.append(rewards)
all_rewards = rewards_list ```

[[10, 20, 30], [10, 20, 30]]

You are training a model using TensorFlow/Keras. Which callback would be MOST useful for preventing overfitting during training?

EarlyStopping

Which preprocessing technique is best suited for handling categorical features with a large number of unique values?

One-hot encoding

What is the primary purpose of using a discount factor (gamma) in reinforcement learning?

To prioritize immediate rewards over long-term rewards.

Consider the following Python code using scikit-learn. What does the code achieve? ```python
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42) ```

It splits the data into training and testing sets.

You're working with a highly imbalanced dataset for a classification task. Which technique would be MOST effective in improving model performance?

Applying resampling techniques (oversampling/undersampling).

What will the following code print? ```python
import numpy as np
import tensorflow as tf
model = tf.keras.Sequential([tf.keras.layers.Dense(1, input_shape=(1,))])
obs = np.array([1.0])
with tf.GradientTape() as tape:
 prediction = model(obs[np.newaxis])
print(prediction.numpy().shape) ```

(1, 1)

